

Machine Learning – Regime Switching HMM

Phase II

November 2021

By: Corbin Apple & Young Seok Seo

Table of Contents




1. Project Summary
2. Model Overview
3. Code Breakdown
4. Data Analysis
5. Empirical Results
6. Next Steps
7. References

Project Summary



Summary

- **Main Idea:** Take advantage of changing market dynamics
- **Model:** Hidden Markov Model, identifying stable, low variance regimes and unstable, high variance regimes
- **Idealized Results:** Construct a replicating portfolio that avoids market drawdowns, but capitalizes on positive performance





Overview

■ Our Idea:

1. Fit a Hidden Markov Model using a rolling window
2. Estimate the current state we are in at each time step, using only historical data
3. Construct a portfolio in which we are fully invested if we are currently in a positive market, and fully divested in a negative market

■ Reasoning:

- Predicting future daily returns is very challenging
- This method allows for model construction purely on actual returns
- Ease of implementation and reproducibility



Breakdown (1 of 11)

- Import necessary packages

```
1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt
4 import matplotlib.dates as mdates
5 import seaborn as sns
6 import datetime
7 from hmmlearn import hmm
8
9 import warnings
10 warnings.filterwarnings("ignore")
```

Code Breakdown (cont.)



Breakdown (2 of 11)

- Import Bank of America historical data

```
1 input_df = pd.read_csv('~\\Desktop\\Corbin SBU\\AMS 520\\Project\\BofA Projects Data\\EOD_20210908.csv',
2                       header = None,
3                       names = [ 'Ticker', # Label columns
4                               'Date',
5                               'Open',
6                               'High',
7                               'Low',
8                               'Close',
9                               'Volume',
10                              'Dividend',
11                              'Stock_split',
12                              'Adj_open',
13                              'Adj_high',
14                              'Adj_low',
15                              'Adj_close',
16                              'Adj_volume' ])
```

Code Breakdown (cont.)



Breakdown (3 of 11)

- Select chosen index (SPY)
- Calculate daily percentage returns
- Calculate volatility of daily returns over selected time window (500 days)

```
1 # Proposed Idea: Create a HMM for the recent Neff days, and for all days after the Neff'th day
2 # Predict which state we are currently in based on the Neff recent days
3 # Be fully invested if in a positive market, fully divested in a negative market
4
5 data = input_df.loc[input_df['Ticker'] == 'SPY'] # Select which index to use for analysis
6 data.reset_index(inplace=True, drop=True)
7
8 Neff = 500 # Lookback length
9
10 # Calculate daily percentage returns
11 Return = 100*(data['Adj_close'] - data['Adj_close'].shift(1)) / data['Adj_close'].shift(1)
12
13 data['Return'] = Return
```

```
1 # Volatility = standard deviation for trailing Neff days
2 Volatility = np.concatenate((np.zeros(Neff),
3                               [np.sqrt(sum((data['Return'][i:Neff+i] -
4                                              np.mean(data['Return'][i:Neff+i])**2) / Neff)
5                               for i in range(len(data)-Neff))])
6
7 data['Volatility'] = Volatility
```

Data Analysis



Data Details

	Ticker	Date	Open	High	Low	Close	Volume	Dividend	Stock_split	Adj_open	Adj_high	Adj_low	Adj_close	Adj_volume	Return	Volatility
0	SPY	1993-01-29	43.9687	43.9687	43.7500	43.9375	1003200.0	0.0	1.0	25.804239	25.804239	25.675889	25.785928	1003200.0	NaN	0.000000
1	SPY	1993-02-01	43.9687	44.2500	43.9687	44.2500	480500.0	0.0	1.0	25.804239	25.969328	25.804239	25.969328	480500.0	0.711238	0.000000
2	SPY	1993-02-02	44.2187	44.3750	44.1250	44.3437	201300.0	0.0	1.0	25.950958	26.042687	25.895968	26.024318	201300.0	0.211751	0.000000
3	SPY	1993-02-03	44.4062	44.8437	44.3750	44.8125	529400.0	0.0	1.0	26.060998	26.317757	26.042687	26.299446	529400.0	1.057196	0.000000
4	SPY	1993-02-04	44.9687	45.0937	44.4687	45.0000	531500.0	0.0	1.0	26.391117	26.464476	26.097678	26.409486	531500.0	0.418410	0.000000
...
7199	SPY	2021-08-31	452.1300	452.4900	450.9200	451.5600	58631140.0	0.0	1.0	452.130000	452.490000	450.920000	451.560000	58631140.0	-0.148155	1.577472
7200	SPY	2021-09-01	452.5600	453.1100	451.5450	451.8000	48667698.0	0.0	1.0	452.560000	453.110000	451.545000	451.800000	48667698.0	0.053149	1.577512
7201	SPY	2021-09-02	453.3200	454.0500	451.9100	453.1900	42479834.0	0.0	1.0	453.320000	454.050000	451.910000	453.190000	42479834.0	0.307658	1.577511
7202	SPY	2021-09-03	451.9800	453.6300	451.5500	453.0800	47155405.0	0.0	1.0	451.980000	453.630000	451.550000	453.080000	47155405.0	-0.024272	1.577528
7203	SPY	2021-09-07	452.7100	452.8100	450.7423	451.4600	51477698.0	0.0	1.0	452.710000	452.810000	450.742300	451.460000	51477698.0	-0.357553	1.577303

7204 rows × 16 columns

Code Breakdown (cont.)



Breakdown (4 of 11)

- Fit the Hidden Markov Models and predict the regime given the current observations

```
1 # Initialize a HMM
2 states = 2
3 max_iterations = 100 # For EM algorithm
4 current_state = np.array([]) # Initialize an array to track the current regimes
5
6 # Exclude first (to calculate trailing volatility) and
7 # second (to have a full set of observations to fit a HMM) Neff observations
8 for i in range(1, len(data) - 2*Neff):
9     # Initialize a Gaussian HMM
10    model = hmm.GaussianHMM(n_components = states, covariance_type="full", n_iter = max_iterations);
11    # Pull Neff observations of Volatility and Returns
12    observations = np.stack((data['Volatility'][Neff+i:Neff*2+i], data['Return'][Neff+i:Neff*2+i]), axis=1)
13    model.fit(observations) # Fit the model to the observations
14    print(f'i = {i}') # Print to ensure loop is running
15    # Model randomly allocates a '0' or a '1' to a state, so check which state has the higher mean return
16    if model.means_[0,1] > model.means_[1,1]:
17        positive_state = 0
18    else:
19        positive_state = 1
20    predictions = model.predict(observations) # Predict the state for each observation
21    if positive_state == 1: # If the state with the higher mean return is state 1, do nothing, if not...
22        pass
23    else: # Switch the predicted regimes to ensure state 1 is always the regime with a greater mean return
24        zeros = np.where(predictions == 0)
25        ones = np.where(predictions == 1)
26        predictions=zeros = 1
27        predictions[ones] = 0
28    # Append the current state to the regime tracker
29    current_state = np.append(current_state,predictions[-1])
30    # Switch values from type float to type int for later calculations
31    current_state = current_state.astype(np.int64)
```

Code Breakdown (cont.)



Breakdown (5 of 11)

- Initial review of the regime switching frequency

```
1 total_regime_switches = sum(np.abs(current_state[1:] - current_state[:-1]))
2 years_of_data = len(current_state) / 252 # Approximately 252 trading days in a year
3 avg_regime_changes_per_year = total_regime_switches / years_of_data
4
5 print(f'Total regime switches: {total_regime_switches}')
6 print(f'Average number of regime switches per year: {avg_regime_changes_per_year:.01f}')
```

Total regime switches: 349

Average number of regime switches per year: 14.2


Code Breakdown (cont.)



Breakdown (6 of 11)

- Visually consider the projected regime switching

```
1 # Plot regime switches to graphically analyze
2
3 dates = [datetime.datetime.strptime(d, "%Y-%m-%d").date() for d in data.Date[2*Neff + 1:]]
4
5 fig, ax = plt.subplots(figsize=(18,4))
6
7 formatter = mdates.DateFormatter("%Y")
8
9 ax.xaxis.set_major_formatter(formatter)
10
11 fmt_half_year = mdates.MonthLocator(interval=24)
12 ax.xaxis.set_major_locator(fmt_half_year)
13
14 ax.plot(dates, current_state, '.', color = 'royalblue');
```



Code Breakdown (cont.)



Breakdown (7 of 11)

- Numerically analyze the mean return and volatility of the two states

```
1 state_0 = np.where(current_state == 0)[0] + 2*Neff + 1
2 state_1 = np.where(current_state == 1)[0] + 2*Neff + 1
```

```
1 # Calculate and view statistics of the two states
2
3 print(f'Number of occurrences of state 0: {len(state_0)}')
4 print(f'Mean return of state 0: {data.loc[state_0].Return.mean():.03f}')
5 print(f'Volatility of state 0: {data.loc[state_0].Return.std():.03f}')
6 print('\n')
7 print(f'Number of occurrences of state 1: {len(state_1)}')
8 print(f'Mean return of state 1: {data.loc[state_1].Return.mean():.03f}')
9 print(f'Volatility of state 1: {data.loc[state_1].Return.std():.03f}')
```

```
Number of occurrences of state 0: 2708
Mean return of state 0: 0.050
Volatility of state 0: 1.469
```

```
Number of occurrences of state 1: 3495
Mean return of state 1: 0.038
Volatility of state 1: 1.040
```

Code Breakdown (cont.)



Breakdown (8 of 11)

- Plot the HMM portfolio against the actual results

```
1 # Calculate the growth of a theoretical portfolio
2 # Assume fully invested if state 1, fully divested if state 0
3
4 current_portfolio = data.Adj_close[2*Neff+1] * np.cumprod(current_state * data.Return[2*Neff+1:] / 100 + 1)
```

```
1 # Plot the return portfolio dynamics
2
3 plot = sns.relplot(x = range(0, len(current_state)),
4                   y = "Adj_close",
5                   data = data[2*Neff+1:],
6                   hue = current_state,
7                   linewidth = 0,
8                   palette = "Set2",
9                   s = 10);
10
11 plt.plot(range(0, len(current_state)), current_portfolio, color='royalblue')
12
13 plot.fig.set_size_inches(18, 10)
```

Code Breakdown (cont.)



Breakdown (9 of 11)




Code Breakdown (cont.)



Breakdown (10 of 11)

- Consider alternative rolling windows

```
1 # Import previous iterations using different Neff window lengths
2 test_portfolio_750 = np.loadtxt("test_portfolio_750.csv")
3 test_portfolio_500 = np.loadtxt("test_portfolio_500.csv")
4 test_portfolio_250 = np.loadtxt("test_portfolio_250.csv")
5
6 fig, ax = plt.subplots(figsize=(18,4))
7
8 ax.plot(range(0,len(current_state)+500), data.Adj_close[2*250+1:], color='slategray');
9 ax.plot(range(1000,len(current_state)+500), test_portfolio_750, color='royalblue');
10 ax.plot(range(500,len(current_state)+500), test_portfolio_500, color='darkgreen');
11 ax.plot(range(0,len(current_state)+500), test_portfolio_250, color='darkviolet');
12
13 ax.legend(['SPY', 'HMM (750-day rolling window)', 'HMM (500-day rolling window)', 'HMM (250-day rolling window)']);
```



Code Breakdown (cont.)



Breakdown (11 of 11)

■ Calculate empirical results

```
1  # Empirical results
2
3  spy_sharpe = data.Return[2*Neff+1:].mean() / data.Return[2*Neff+1:].std()
4
5  test_portfolio_750_returns = (test_portfolio_750[1:] - test_portfolio_750[:-1]) / test_portfolio_750[:-1]
6  test_portfolio_750_sharpe = test_portfolio_750_returns.mean() / test_portfolio_750_returns.std()
7
8  test_portfolio_500_returns = (test_portfolio_500[1:] - test_portfolio_500[:-1]) / test_portfolio_500[:-1]
9  test_portfolio_500_sharpe = test_portfolio_500_returns.mean() / test_portfolio_500_returns.std()
10
11 test_portfolio_250_returns = (test_portfolio_250[1:] - test_portfolio_250[:-1]) / test_portfolio_250[:-1]
12 test_portfolio_250_sharpe = test_portfolio_250_returns.mean() / test_portfolio_250_returns.std()
13
14 print(f'SPY average return: {spy_mean_return:.03f}')
15 print(f'SPY volatility: {spy_vol:.03f}')
16 print(f'SPY Sharpe Ratio: {spy_sharpe:.03f}')
17 print('\n')
18 print(f'HMM (750-day rolling window) average return: {data.Return[2*Neff+1:].mean():.05f}')
19 print(f'HMM (750-day rolling window) volatility: {data.Return[2*Neff+1:].std():.05f}')
20 print(f'HMM (750-day rolling window) Sharpe Ratio: {test_portfolio_750_sharpe:.03f}')
21 print('\n')
22 print(f'HMM (500-day rolling window) average return: {test_portfolio_500_returns.mean():.05f}')
23 print(f'HMM (500-day rolling window) volatility: {test_portfolio_500_returns.std():.05f}')
24 print(f'HMM (500-day rolling window) Sharpe Ratio: {test_portfolio_500_sharpe:.03f}')
25 print('\n')
26 print(f'HMM (250-day rolling window) average return: {test_portfolio_250_returns.mean():.05f}')
27 print(f'HMM (250-day rolling window) volatility: {test_portfolio_250_returns.std():.05f}')
28 print(f'HMM (250-day rolling window) Sharpe Ratio: {test_portfolio_250_sharpe:.03f}')
```


Empirical Results



Initial Results

SPY average return: 0.043
SPY volatility: 1.246
SPY Sharpe Ratio: 0.035

HMM (750-day rolling window) average return: 0.04337
HMM (750-day rolling window) volatility: 1.24552
HMM (750-day rolling window) Sharpe Ratio: 0.026

HMM (500-day rolling window) average return: 0.00021
HMM (500-day rolling window) volatility: 0.00781
HMM (500-day rolling window) Sharpe Ratio: 0.027

HMM (250-day rolling window) average return: 0.00019
HMM (250-day rolling window) volatility: 0.00747
HMM (250-day rolling window) Sharpe Ratio: 0.025

Sharpe Ratios: SPY > HMM (500) > HMM (750) > HMM (250)



Future Ideas for Phase III

- Performing this same analysis with adjusted factors, including using:
 - Different rolling window lengths (e.g. 10-1,000 days), in search of an “optimal” lookback timeframe
 - Other indices (domestic and international), to see if an “optimal” window length is consistent across funds and markets
 - Incorporating transaction costs each time we switch regimes, with the goal of penalizing models that recommend frequent rebalancings
- Should we decide to go down the route of predicting day-ahead returns, incorporating these forecasts into our Hidden Markov Model, we will likely try to implement a standard time series model (e.g. GARCH, ARMA-GARCH)
 - Predict future states by fitting forecasted observations
 - Make optimal investment decisions in advance of realized returns
- Reorganize main code as a single function to support easy reproducibility of results
- Produce results with visualization package (based on Shiny framework, QuantStats)

References



Reference List

- Peter Nystrup, Henrik Madsen & Erik Lindström (2018) Dynamic portfolio optimization across hidden market regimes, *Quantitative Finance*, 18:1, 83-95, DOI: 10.1080/14697688.2017.1342857
- Wang M, Lin Y-H, Mikhelson I. Regime-Switching Factor Investing with Hidden Markov Models. *Journal of Risk and Financial Management*. 2020; 13(12):311. <https://doi.org/10.3390/jrfm13120311>